

Telco Churn Analysis

Import dependencies

```
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
library(scales)
library(ggplot2)
library(corrplot)

## corrplot 0.84 loaded
library(gridExtra)

##
## Attaching package: 'gridExtra'
## The following object is masked from 'package:dplyr':
##
##   combine
library(ggthemes)
library(caret)

## Loading required package: lattice
library(MASS)

##
## Attaching package: 'MASS'
## The following object is masked from 'package:dplyr':
##
##   select
library(randomForest)

## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:gridExtra':
##
##      combine
## The following object is masked from 'package:ggplot2':
##
##      margin
## The following object is masked from 'package:dplyr':
##
##      combine
library(party)
```

```
## Loading required package: grid
## Loading required package: mvtnorm
## Loading required package: modeltools
## Loading required package: stats4
## Loading required package: strucchange
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
## Loading required package: sandwich
```

```
library(MLmetrics)

##
## Attaching package: 'MLmetrics'
## The following objects are masked from 'package:caret':
##
##      MAE, RMSE
## The following object is masked from 'package:base':
##
##      Recall
library(rpart)
library(rpart.plot)
library(precrec)
```


Read in data and preprocess

```
telcoChurn <- read.csv(file = 'telco_churn_edited.csv')
telcoChurn = subset(telcoChurn, select=-c(customerID))
telcoChurn$Churn <- as.factor(telcoChurn$Churn)
telcoChurn <- na.omit(telcoChurn)
```

Exploratory Data Analysis


Tenure against Churn

```
tenureCounts <- table(telcoChurn$tenure)
barplot(
  tenureCounts,
  main="Tenure Distribution",
  xlab="Tenure",
  col="skyblue",
  border=F
)
```




```
# percentage exited for each tenure
tenure <- telcoChurn %>% count(Churn, tenure)
tenure$Churn <- as.factor(tenure$Churn)
tenure$Churn <- ifelse(tenure$Churn == "No", "No Churn", "Churned")

ggplot(tenure, aes(fill=Churn, y=n, x=tenure))+
  geom_bar(position="fill", stat="identity")+
  scale_fill_manual(values = c("skyblue", scales::alpha("red", .5)))
```



```
# boxplot for tenure against churn
boxplot(
  tenure ~ Churn,
  data=telcoChurn,
  main="Tenure against Churn",
  xlab="Churned",
  ylab="Tenure"
)
```


Tenure against Churn



Contract Type against Churn


```
ContractType <- table(telcoChurn$Contract)
barplot(
  ContractType,
  main="Contract Type Distribution",
  xlab="Contract",
  col="skyblue",
  border=F
)
```

Contract Type Distribution



```
# percentage exited for each contract type
Contract <- telcoChurn %>% count(Churn, Contract)
Contract$Churn <- as.factor(Contract$Churn)
Contract$Churn <- ifelse(Contract$Churn == "No", "No Churn", "Churned")

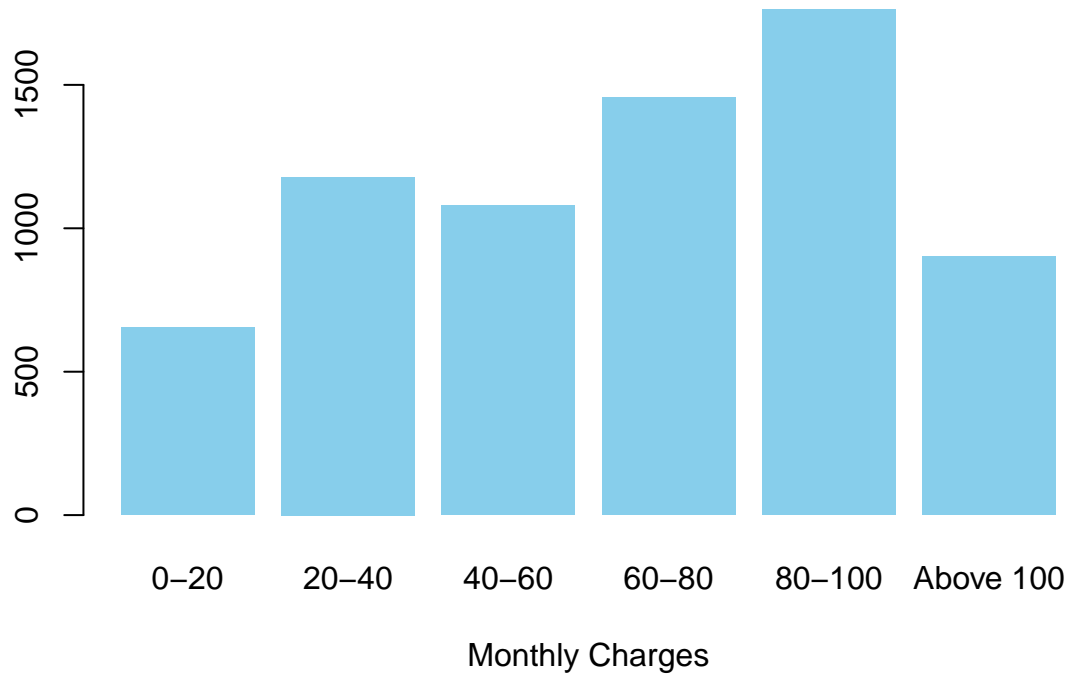
ggplot(Contract, aes(fill=Churn, y=n, x=Contract))+
  geom_bar(position="fill", stat="identity")+
  scale_fill_manual(values = c("skyblue", scales::alpha("red", .5)))
```



Monthly Charges against Churn


```
MonthlyChargesCounts <- table(telcoChurn$MonthlyChargesIntervals)
barplot(
  MonthlyChargesCounts,
  main="Monthly Charges Distribution",
  xlab="Monthly Charges",
  col="skyblue",
  border=F
)
```

Monthly Charges Distribution



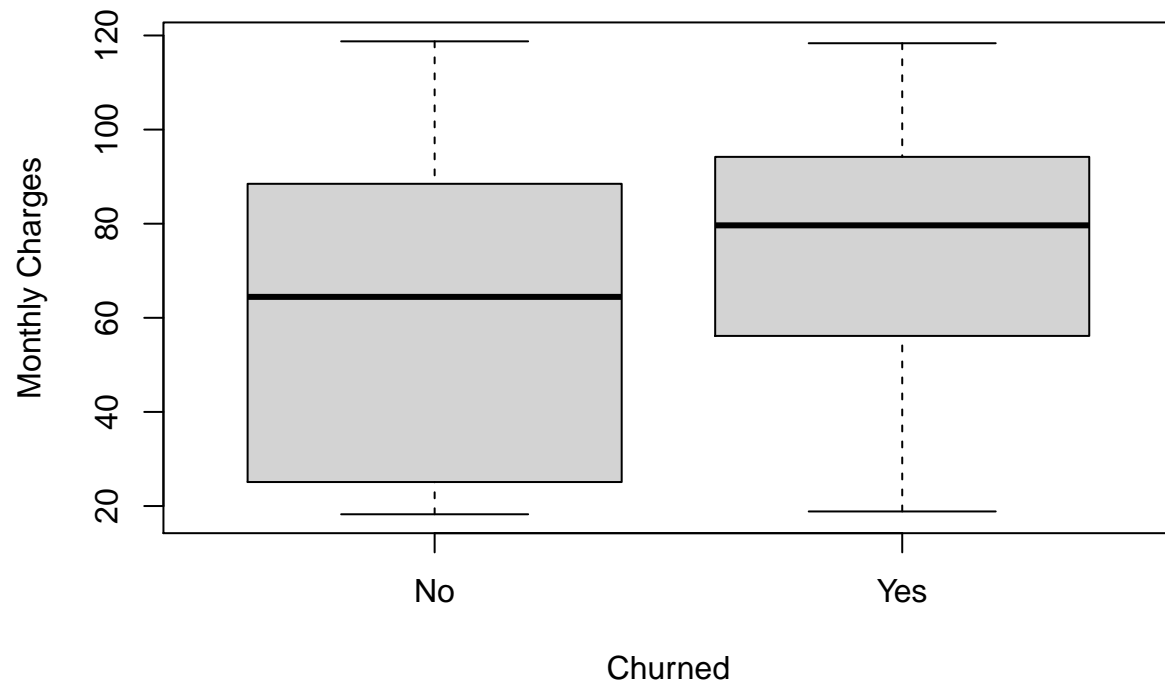
```
# percentage exited for Montly Charge Range
MonthlyChargesIntervals <- telcoChurn %>% count(Churn, MonthlyChargesIntervals)
MonthlyChargesIntervals$Churn <- as.factor(MonthlyChargesIntervals$Churn)
MonthlyChargesIntervals$Churn <- ifelse(MonthlyChargesIntervals$Churn == "No", "No Churn", "Churned")

ggplot(MonthlyChargesIntervals, aes(fill=Churn, y=n, x=MonthlyChargesIntervals))+
  geom_bar(position="fill", stat="identity")+
  scale_fill_manual(values = c("skyblue", scales::alpha("red", .5)))
```

```
# boxplot for monthly charges against churn
boxplot(
  MonthlyCharges ~ Churn,
  data=telcoChurn,
  main="Monthly Charges against Churned",
  xlab="Churned",
  ylab="Monthly Charges"
)
```

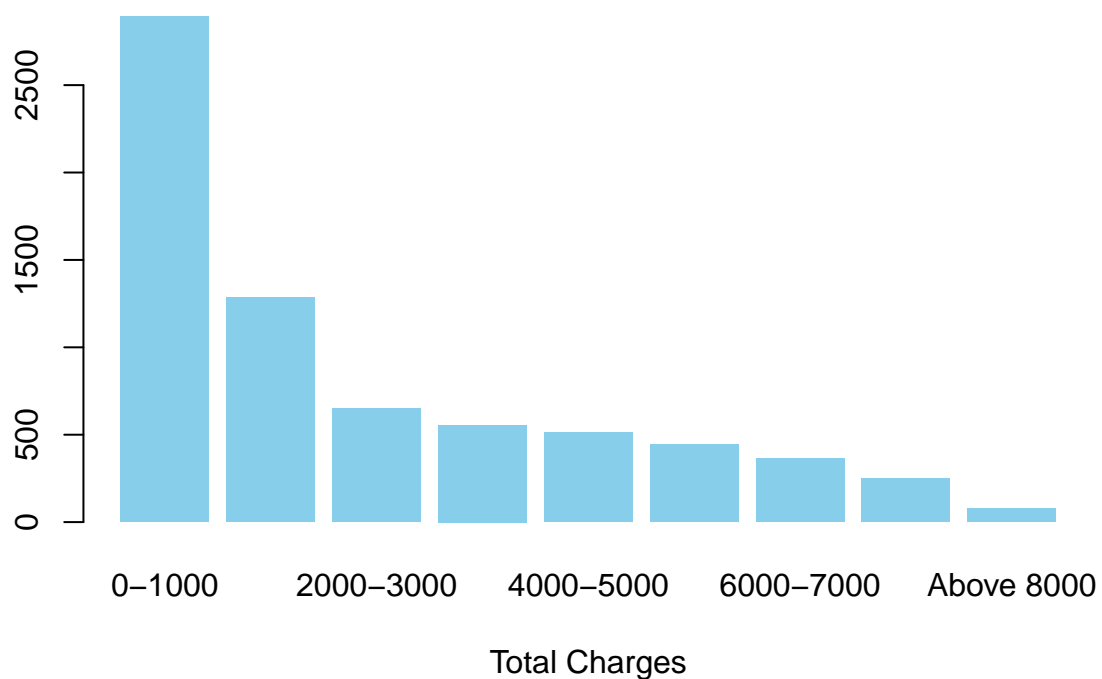
Monthly Charges against Churned



Total Charges against Churn


```
TotalChargesCounts <- table(telcoChurn$TotalChargesIntervals)
barplot(
  TotalChargesCounts,
  main="Total Charges Distribution",
  xlab="Total Charges",
  col="skyblue",
  border=F
)
```

Total Charges Distribution



```
# percentage exited for Montly Charge Range
TotalChargesIntervals <- telcoChurn %>% count(Churn, TotalChargesIntervals)
TotalChargesIntervals$Churn <- as.factor(TotalChargesIntervals$Churn)
TotalChargesIntervals$Churn <- ifelse(TotalChargesIntervals$Churn == "No", "No Churn", "Churned")

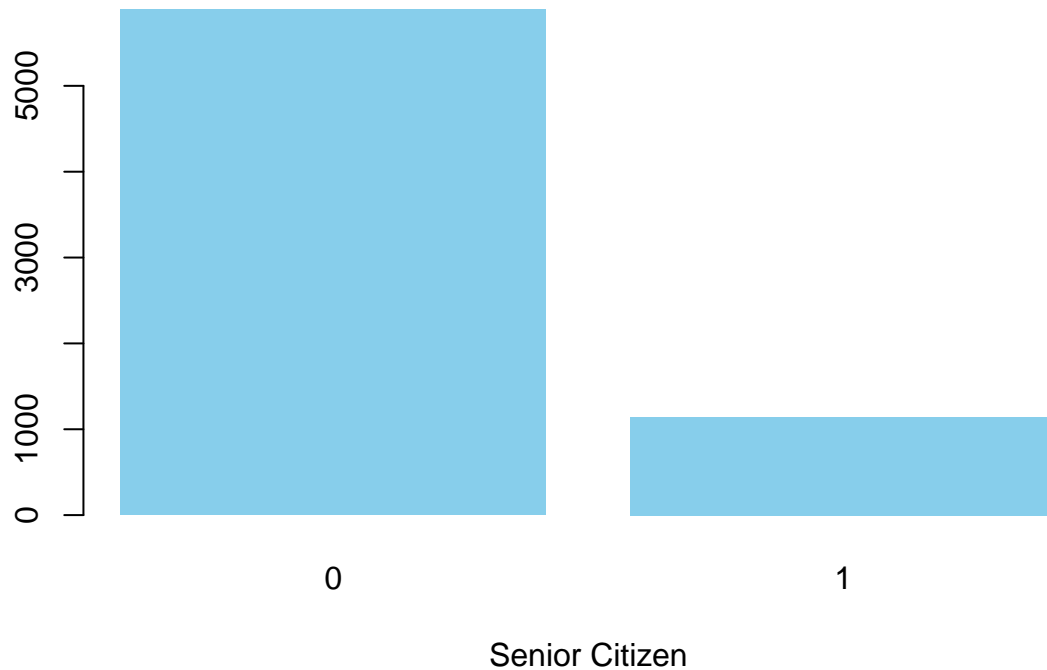
ggplot(TotalChargesIntervals, aes(fill=Churn, y=n, x=TotalChargesIntervals))+
  geom_bar(position="fill", stat="identity")+
  scale_fill_manual(values = c("skyblue", scales::alpha("red", .5)))
```



Senior Citizen against Churn


```
SeniorCitizenscount <- table(telcoChurn$SeniorCitizen)
barplot(
  SeniorCitizenscount,
  main="Senior Citizen Ratio",
  xlab="Senior Citizen",
  col="skyblue",
  border=F
)
```

Senior Citizen Ratio



```
# percentage exited for Montly Charge Range
SeniorCitizen <- telcoChurn %>% count(Churn, SeniorCitizen)
SeniorCitizen$Churn <- as.factor(SeniorCitizen$Churn)
SeniorCitizen$Churn <- ifelse(SeniorCitizen$Churn == "No", "No Churn", "Churned")

ggplot(SeniorCitizen, aes(fill=Churn, y=n, x=SeniorCitizen))+
  geom_bar(position="fill", stat="identity")+
  scale_fill_manual(values = c("skyblue", scales::alpha("red", .5)))
```



Training

Train-test-split

```
# train test split
idx = createDataPartition(telcoChurn$Churn, p=0.7, list=FALSE)
set.seed(42)
train = telcoChurn[idx,]
test = telcoChurn[-idx,]

train$Churn = ifelse(train$Churn == "Yes",1,0)
test$Churn = ifelse(test$Churn == "Yes",1,0)
train$Churn = as.factor(train$Churn)
test$Churn = as.factor(test$Churn)

dim(train); dim(test)

## [1] 4924 22
## [1] 2108 22
```

Logistic Regression

```
# logistic regression to predict churn
logreg = glm(Churn ~ .,
             family=binomial(link="logit"),
```

```

data=train)

print(summary(logreg))

##
## Call:
## glm(formula = Churn ~ ., family = binomial(link = "logit"), data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9698  -0.6763  -0.3052   0.6772   3.1916
##
## Coefficients: (7 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    1.0749035   1.0268513    1.047 0.295194
## genderMale     -0.0463684   0.0774897   -0.598 0.549586
## SeniorCitizen    0.1354064   0.1018150    1.330 0.183542
## PartnerYes      0.0422181   0.0923358    0.457 0.647510
## DependentsYes   -0.0805063   0.1070295   -0.752 0.451938
## tenure         -0.0363180   0.0092040   -3.946 7.95e-05 ***
## PhoneServiceYes  0.0265474   0.7995017    0.033 0.973511
## MultipleLinesNo phone service      NA         NA         NA      NA
## MultipleLinesYes  0.5757736   0.2129948    2.703 0.006867 **
## InternetServiceFiber optic    1.9189455   0.9632208    1.992 0.046347 *
## InternetServiceNo -1.6258190   1.0090429   -1.611 0.107126
## OnlineSecurityNo internet service      NA         NA         NA      NA
## OnlineSecurityYes -0.1485496   0.2146207   -0.692 0.488843
## OnlineBackupNo internet service      NA         NA         NA      NA
## OnlineBackupYes   0.0158870   0.2101957    0.076 0.939752
## DeviceProtectionNo internet service      NA         NA         NA      NA
## DeviceProtectionYes  0.1450105   0.2108702    0.688 0.491656
## TechSupportNo internet service      NA         NA         NA      NA
## TechSupportYes    -0.0978890   0.2158129   -0.454 0.650129
## StreamingTVNo internet service      NA         NA         NA      NA
## StreamingTVYes     0.6220082   0.3939108    1.579 0.114323
## StreamingMoviesNo internet service      NA         NA         NA      NA
## StreamingMoviesYes  0.5565557   0.3912933    1.422 0.154925
## ContractOne year  -0.5344155   0.1282571   -4.167 3.09e-05 ***
## ContractTwo year  -1.1114079   0.2038995   -5.451 5.02e-08 ***
## PaperlessBillingYes  0.3707521   0.0897391    4.131 3.60e-05 ***
## PaymentMethodCredit card (automatic) -0.1003492   0.1357758   -0.739 0.459858
## PaymentMethodElectronic check    0.2993887   0.1135679    2.636 0.008384 **
## PaymentMethodMailed check -0.1238955   0.1377694   -0.899 0.368495
## MonthlyCharges    -0.0352845   0.0397288   -0.888 0.374468
## MonthlyChargesIntervals20-40 -0.1583337   0.2487169   -0.637 0.524384
## MonthlyChargesIntervals40-60  0.2831188   0.3976267    0.712 0.476451
## MonthlyChargesIntervals60-80 -0.0087914   0.5125293   -0.017 0.986315
## MonthlyChargesIntervals80-100 -0.0399477   0.6116405   -0.065 0.947925
## MonthlyChargesIntervalsAbove 100  0.2140849   0.7088860    0.302 0.762651
## TotalCharges      -0.0007546   0.0002023   -3.730 0.000192 ***
## TotalChargesIntervals1000-2000  0.4057784   0.1914898    2.119 0.034085 *
## TotalChargesIntervals2000-3000  1.3437144   0.3238869    4.149 3.34e-05 ***
## TotalChargesIntervals3000-4000  2.2043882   0.4640929    4.750 2.04e-06 ***
## TotalChargesIntervals4000-5000  2.9534158   0.6147289    4.804 1.55e-06 ***

```

```
## TotalChargesIntervals5000-6000      3.9029024  0.7637289   5.110 3.22e-07 ***
## TotalChargesIntervals6000-7000      4.8475966  0.9177026   5.282 1.28e-07 ***
## TotalChargesIntervals7000-8000      5.5574340  1.0701044   5.193 2.07e-07 ***
## TotalChargesIntervalsAbove 8000      5.6977092  1.3302252   4.283 1.84e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 5702.8  on 4923  degrees of freedom
## Residual deviance: 4110.9  on 4887  degrees of freedom
## AIC: 4184.9
##
## Number of Fisher Scoring iterations: 6
```

Feature importance using deviance

```
# feature importance: the steeper the drop in deviance the more important the feature
anova(logreg, test="Chisq")
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Churn
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                                4923     5702.8
## gender              1      0.75     4922     5702.0 0.385995
## SeniorCitizen       1     82.82     4921     5619.2 < 2.2e-16 ***
## Partner             1    109.76     4920     5509.4 < 2.2e-16 ***
## Dependents         1     29.51     4919     5479.9 5.563e-08 ***
## tenure             1    591.08     4918     4888.8 < 2.2e-16 ***
## PhoneService       1      0.28     4917     4888.6 0.594329
## MultipleLines       1    128.27     4916     4760.3 < 2.2e-16 ***
## InternetService     2    427.35     4914     4332.9 < 2.2e-16 ***
## OnlineSecurity      1     26.00     4913     4306.9 3.415e-07 ***
## OnlineBackup        1      1.11     4912     4305.8 0.292842
## DeviceProtection    1      0.09     4911     4305.7 0.759166
## TechSupport         1     17.68     4910     4288.1 2.613e-05 ***
## StreamingTV         1     19.62     4909     4268.4 9.449e-06 ***
## StreamingMovies     1      6.36     4908     4262.1 0.011643 *
## Contract           2     44.95     4906     4217.1 1.732e-10 ***
## PaperlessBilling    1     18.47     4905     4198.7 1.728e-05 ***
## PaymentMethod       3     23.04     4902     4175.6 3.970e-05 ***
## MonthlyCharges      1      0.76     4901     4174.9 0.383383
## MonthlyChargesIntervals 5     18.16     4896     4156.7 0.002756 **
## TotalCharges        1      9.05     4895     4147.7 0.002628 **
## TotalChargesIntervals 8     36.77     4887     4110.9 1.267e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


Accuracy

```
# evaluating logistic regression model against test data
logreg.scores <- predict(logreg,newdata=test,type='response')

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading

logreg.pred <- ifelse(logreg.scores > 0.5,1,0)
misclassError <- mean(logreg.pred != test$Churn)
print(paste('Logistic Regression Accuracy',1-misclassError))

## [1] "Logistic Regression Accuracy 0.805028462998102"
```

Confusion Matrix

```
print("Confusion Matrix for Logistic Regression"); table(Predicted = logreg.pred, Actual = test$Churn)

## [1] "Confusion Matrix for Logistic Regression"



##           Actual
## Predicted    0    1
##           0 1398  261
##           1   150  299
```

ROC, PRC curves

```
precrec.logreg <- evalmod(scores = logreg.scores, labels = test$Churn)
print(precrec.logreg)


##
##      === AUCs ===
##
##      Model name Dataset ID Curve type      AUC
##      1          m1      1      ROC 0.8638774
##      2          m1      1      PRC 0.6755818
##
##
##      === Input data ===
##
##      Model name Dataset ID # of negatives # of positives
##      1          m1      1          1548          560

autoplot(precrec.logreg)
```



Decision Tree

```
cart <- rpart(Churn ~ .,  
              data=train,  
              method = 'class',  
              control=rpart.control(minsplit = 2, cp = 0)  
            )  
plotcp(cart)
```



Optimisation of CP value

```
CError.cap <- cart$cptable[which.min(cart$cptable[, "xerror"]), "xerror"] + cart$cptable[which.min(cart$cptable[, "xerror"]), "xerror"]

# Find the optimal CP region whose CV error is just below CError.cap in maximal tree cart1.
i <- 1; j <- 4
while (cart$cptable[i, j] > CError.cap) {
  i <- i + 1
}

# Get geometric mean of the two identified CP values in the optimal region if optimal tree has at least 2 nodes
cp.opt = ifelse(i > 1, sqrt(cart$cptable[i, 1] * cart$cptable[i-1, 1]), 1)

cart.opt <- prune(cart, cp = cp.opt)
cart.opt$variable.importance
```

	Contract	tenure	TotalCharges	Intervals
##	291.35285121	271.05518868	166.01119860	
##	OnlineSecurity	TotalCharges	TechSupport	
##	142.37651359	130.80241298	125.73512692	
##	DeviceProtection	MonthlyCharges	InternetService	
##	108.65083704	105.08046232	99.09262490	
##	MonthlyChargesIntervals	MultipleLines	StreamingTV	
##	90.43297869	46.27607262	16.63781288	
##	OnlineBackup	PaymentMethod	PaperlessBilling	
##	16.60431846	5.03950863	2.28343428	
##	Dependents	Partner		
##	0.08954644	0.04477322		

Accuracy

```
tree.pred <- predict(cart.opt, test, type="class")
tree.pred.scores <- predict(cart.opt, test, type="prob")
table.pred <- table(Predicted = tree.pred, Actual = test$Churn)
print(paste('Decision Tree Accuracy', sum(diag(table.pred))/sum(table.pred)))
```

```
## [1] "Decision Tree Accuracy 0.790322580645161"
```

Confusion Matrix

```
# Note: accuracy is 80% because of unbalanced dataset; most data points have Churn = 0. From the confus
print("Confusion Matrix for Decision Tree"); table(Predicted = tree.pred, Actual = test$Churn)
```


```
## [1] "Confusion Matrix for Decision Tree"
```

```
##           Actual
## Predicted    0    1
##           0 1417  311
##           1  131  249
```

ROC, PRC curves

```
precrec.tree <- evalmod(scores = tree.pred.scores[, 2], labels = test$Churn)
print(precrec.tree)
```

```
##
##      === AUCs ===
##
##      Model name Dataset ID Curve type      AUC
##      1         m1         1      ROC 0.8301483
##      2         m1         1      PRC 0.6039751
##
##
##      === Input data ===
##
##      Model name Dataset ID # of negatives # of positives
##      1         m1         1         1548         560
autoplots(precrec.tree)
```



Random Forest


```
rfmodel <- randomForest(Churn ~ ., data = train, proximity = TRUE)
```

```
rfmodel
```

```
##
## Call:
## randomForest(formula = Churn ~ ., data = train, proximity = TRUE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 20.53%
## Confusion matrix:
##      0   1 class.error
## 0 3289 326  0.09017981
## 1  685 624  0.52330023
```

OOB error plot for initial model

```
# model based on err.rate matrix: [OOB, No, Yes]
oob.error.data <- data.frame(Trees=rep(1:nrow(rfmodel$err.rate), times=3), Type=rep(c("OOB", "0", "1"),
ggplot(data=oob.error.data, aes(x=Trees, y=Error))+geom_line(aes(color=Type))
```



Accuracy for initial model

```
rfmodel.pred <- predict(rfmodel, test, type="response")
rfmodel.pred.scores <- predict(rfmodel, test, type="prob")
table.rfmodel.pred <- table(Predicted = rfmodel.pred, Actual = test$Churn)
print(paste('Random Forest Accuracy', sum(diag(table.rfmodel.pred))/sum(table.rfmodel.pred)))

## [1] "Random Forest Accuracy 0.799810246679317"
```

Confusion Matrix for initial model

```
print("Confusion Matrix for Random Forest"); table(Predicted = rfmodel.pred, Actual = test$Churn)

## [1] "Confusion Matrix for Random Forest"

##           Actual
## Predicted    0    1
##           0 1410  284
##           1  138  276
```

ROC, PRC curves for initial model

```
precrec.rf <- evalmod(scores = rfmodel.pred.scores[, 2], labels = test$Churn)
print(precrec.rf)

##
## === AUCs ===
```

```
##
##      Model name Dataset ID Curve type      AUC
##      1          m1      1      ROC 0.8498823
##      2          m1      1      PRC 0.6640752
##
##
##      === Input data ===
##
##      Model name Dataset ID # of negatives # of positives
##      1          m1      1          1548          560
```

Optimisation of no. of variable in each internal node

```
# optimize no. of variables at each internal node in tree
oob.values <- vector(length=10)
for(i in 1:10){
  temp.model <- randomForest(Churn ~ ., data = telcoChurn, mtry=i, ntree=1000)

  #store OOB error rate for each random forest that uses diff value of i
  oob.values[i] <- temp.model$err.rate[nrow(temp.model$err.rate),1]
}

oob.values

## [1] 0.2127418 0.2006542 0.2037827 0.2025028 0.2044937 0.2044937 0.2070535
## [8] 0.2087600 0.2076223 0.2100398

# no. of variables = 2 gives lowest oob err.rate

rfmodeloptim <- randomForest(Churn ~ ., data = train, mtry=2, proximity = TRUE, type='classification')

rfmodeloptim

##
## Call:
## randomForest(formula = Churn ~ ., data = train, mtry = 2, proximity = TRUE, type = "classification")
##      Type of random forest: classification
##      Number of trees: 500
## No. of variables tried at each split: 2
##
##      OOB estimate of error rate: 20.09%
## Confusion matrix:
##      0  1 class.error
## 0 3315 300  0.08298755
## 1  689 620  0.52635600

# error rate reduced
```

Accuracy for optimised model

```
rfmodeloptim.pred <- predict(rfmodeloptim, test, type="response")
rfmodeloptim.pred.scores <- predict(rfmodeloptim, test, type="prob")
table.rfmodeloptim.pred <- table(Predicted = rfmodeloptim.pred, Actual = test$Churn)
print(paste('Random Forest Accuracy', sum(diag(table.rfmodeloptim.pred))/sum(table.rfmodeloptim.pred)))

## [1] "Random Forest Accuracy 0.796489563567362"
```

Confusion Matrix for optimised model

```
print("Confusion Matrix for Random Forest"); table(Predicted = rfmodeloptim.pred, Actual = test$Churn)

## [1] "Confusion Matrix for Random Forest"


##           Actual
## Predicted    0    1
##           0 1421  302
##           1  127  258
```

ROC, PRC curves for optimised model

```
precrec.rf <- evalmod(scores = rfmodeloptim.pred.scores[, 2], labels = test$Churn)
print(precrec.rf)
```

```
##
##      === AUCs ===
##
##      Model name Dataset ID Curve type      AUC
##      1          m1       1      ROC 0.8497658
##      2          m1       1      PRC 0.6668204
##
##
##      === Input data ===
##
##      Model name Dataset ID # of negatives # of positives
##      1          m1       1          1548           560
```

```
autoplot(precrec.rf)
```




OOB error plot for optimised model

```
oob.error.data <- data.frame(Trees=rep(1:nrow(rfmodel$err.rate), times=3), Type=rep(c("OOB", "0", "1"),
```



```
ggplot(data=oob.error.data, aes(x=Trees, y=Error))+geom_line(aes(color=Type))
```



Feature Importances across models

```
imp.logreg <- varImp(logreg, scale = FALSE)
imp.logreg
```

##	Overall
## genderMale	0.59838111
## SeniorCitizen	1.32992665
## PartnerYes	0.45722380
## DependentsYes	0.75218836
## tenure	3.94587331
## PhoneServiceYes	0.03320498
## MultipleLinesYes	2.70322796
## InternetServiceFiber optic	1.99221781
## InternetServiceNo	1.61124858
## OnlineSecurityYes	0.69214973
## OnlineBackupYes	0.07558182
## DeviceProtectionYes	0.68767678
## TechSupportYes	0.45358284
## StreamingTVYes	1.57905858
## StreamingMoviesYes	1.42234902
## ContractOne year	4.16675335
## ContractTwo year	5.45076283
## PaperlessBillingYes	4.13144246

```
## PaymentMethodCredit card (automatic) 0.73908013
## PaymentMethodElectronic check 2.63620806
## PaymentMethodMailed check 0.89929640
## MonthlyCharges 0.88813459
## MonthlyChargesIntervals20-40 0.63660211
## MonthlyChargesIntervals40-60 0.71202178
## MonthlyChargesIntervals60-80 0.01715294
## MonthlyChargesIntervals80-100 0.06531244
## MonthlyChargesIntervalsAbove 100 0.30200190
## TotalCharges 3.72962419
## TotalChargesIntervals1000-2000 2.11906029
## TotalChargesIntervals2000-3000 4.14871423
## TotalChargesIntervals3000-4000 4.74988570
## TotalChargesIntervals4000-5000 4.80441973
## TotalChargesIntervals5000-6000 5.11032405
## TotalChargesIntervals6000-7000 5.28231786
## TotalChargesIntervals7000-8000 5.19335686
## TotalChargesIntervalsAbove 8000 4.28326648
```

```
imp.tree <- varImp(cart.opt, scale = FALSE)
imp.tree
```

```
## Overall
## Contract 291.352851
## DeviceProtection 18.250475
## InternetService 96.466156
## MonthlyCharges 7.280132
## MultipleLines 2.842111
## OnlineBackup 32.131150
## OnlineSecurity 343.114904
## PaperlessBilling 5.985914
## PaymentMethod 246.841296
## TechSupport 303.520728
## tenure 390.958651
## TotalCharges 106.418423
## TotalChargesIntervals 58.089070
## gender 0.000000
## SeniorCitizen 0.000000
## Partner 0.000000
## Dependents 0.000000
## PhoneService 0.000000
## StreamingTV 0.000000
## StreamingMovies 0.000000
## MonthlyChargesIntervals 0.000000
```

```
imp.rfmodeloptim <- varImp(rfmodeloptim, scale = FALSE)
imp.rfmodeloptim
```

```
## Overall
## gender 22.521053
## SeniorCitizen 20.729294
## Partner 21.668583
## Dependents 20.240538
## tenure 165.620070
## PhoneService 7.233984
```

## MultipleLines	25.219132
## InternetService	46.396192
## OnlineSecurity	59.639262
## OnlineBackup	37.534687
## DeviceProtection	33.560862
## TechSupport	54.440465
## StreamingTV	20.359073
## StreamingMovies	21.690099
## Contract	93.283007
## PaperlessBilling	31.439824
## PaymentMethod	55.068175
## MonthlyCharges	136.733289
## MonthlyChargesIntervals	46.835585
## TotalCharges	161.657314
## TotalChargesIntervals	65.420257